

Fraud Detection in Python

Kai BANTO ECKERT

July 20, 2021

Abstract

A typical organization loses an estimated 5% of its yearly revenue to fraud. In this course, you will learn how to fight fraud by using data. For example, you'll learn how to apply supervised learning algorithms to detect fraudulent behavior similar to past ones, as well as unsupervised learning methods to discover new types of fraud activities. Moreover, in fraud analytics you often deal with highly imbalanced datasets when classifying fraud versus non-fraud, and during this course you will pick up some techniques on how to deal with that. The course provides a mix of technical and theoretical insights and shows you hands-on how to practically implement fraud detection models. In addition, you will get tips and advice from real-life experience to help you prevent making common mistakes in fraud analytics.

1 Introduction and preparing your data

Did you know that a typical organization loses 5% of its revenue to fraud each year? In fact, it is estimated that fraud is costing the UK economy 73 billion British pounds each year. As such, you can say, fraud poses a serious problem to almost all companies. This course teaches you how you can tackle fraud as a data scientist, and thereby make a tangible impact on your company.

1.1 What is fraud?

Fraudulent behavior can be found in many different areas. Credit card fraud is perhaps the most famous example, and also in the insurance industry, fraud is a well-known issue. But it is much more broadly present than that. For example even all e-commerce businesses need to continuously assess whether client transactions on their website are legit. Detecting fraud is typically challenging because of these four characteristics of fraud described here. First of all, fraud cases are in a minority, sometimes only one-hundredth percent of a companies' transactions are fraudulent. Fraudsters will also try their best to "blend" in and conceal their activities. Moreover, fraudsters will find new methods to avoid getting caught, and change their behavior over time. Lastly, fraudsters oftentimes work together and organize their activities in a network, making it harder to detect. It can be that multiple client accounts are involved around one fraud case. Let's illustrate this with an example.

1.2 Fraud detection is challenging

Have you ever played "Where is Waldo" or "Find the odd one out"? Like in the game, in fraud detection you'll need to train an algorithm to pick a well concealed observation out of many normal observations. It looks like the other clovers, but it deviates slightly as it has 3 leaves instead of four. That one was easy, but it does get much harder when we're working with numbers. This is much more like in real life, we'll need to find a fraud case based on numbers. This illustrates a typical fraud detection problem really well: based on data, you'll need to train an algorithm to find the odd one out among many normal observations.

1.3 How companies deal with fraud

As a data scientist working on fraud analytics, you'll often be asked to improve existing fraud detection systems. You'll maybe find that the company already uses a rules based system to filter out strange cases. Or that the fraud analytics team checks the news for suspicious names, or keeps track of external hit lists from the police to reference check against the client base. All these existing methods can be useful for your machine learning model, as you can use them as inputs in your analysis. But do be mindful when using labels that come out of existing rules based systems; you should always ask yourself whether the labels are reliable as they might not catch all fraudulent cases.

1.4 Let's have a look at some data

In this chapter we'll explore a dataset on credit card transactions. We have 29 features available, and a `class` variable, containing information about whether the transaction is fraudulent or not. We have data on 5050 transactions in total. This should be enough for training our first algorithm on. Now let's have a look at this credit card data in more detail!

1.5 Checking the fraud to non-fraud ratio

In this chapter, you will work on `creditcard_sampledata_3.csv`, a dataset containing credit card transactions data. Fraud occurrences are fortunately an **extreme minority** in these transactions.

However, Machine Learning algorithms usually work best when the different classes contained in the dataset are more or less equally present. If there are few cases of fraud, then there's little data to learn how to identify them. This is known as **class imbalance**, and it's one of the main challenges of fraud detection.

1.5 Checking the fraud to non-fraud ratio

Let's explore this dataset, and observe this class imbalance problem.

```
1 | import pandas as pd
2 | import config
3 |
4 | df = pd.read_csv(config.DATAPATH / "chapter_1/creditcard_sampledata.csv")
5 | print(df.head())

1 |      Unnamed: 0  Time      V1      V2  ...      V27      V28  Amount  Class
2 | 0              0    64  1.212511 -0.099054 ...  0.020370  0.017037   34.70      0
3 | 1              1    64 -0.658305  0.406791 ... -0.094192 -0.092493   54.99      0
4 | 2              2   124  1.105253  0.541842 ...  0.000208  0.026167    6.24      0
5 | 3              3   128  1.239495 -0.182609 ...  0.036867  0.010963    8.80      0
6 | 4              4   132 -1.571359  1.687508 ... -0.481570 -0.167897   10.00      0
7 | [5 rows x 32 columns]

1 | df.info()

1 | <class 'pandas.core.frame.DataFrame'>
2 | RangeIndex: 8000 entries, 0 to 7999
3 | Data columns (total 32 columns):
4 |  #   Column      Non-Null Count  Dtype
5 |  ---  ---
6 |  0   Unnamed: 0    8000 non-null     int64
7 |  1   Time          8000 non-null     int64
8 |  2   V1             8000 non-null     float64
9 |  3   V2             8000 non-null     float64
10 |  4   V3             8000 non-null     float64
11 |  5   V4             8000 non-null     float64
12 |  6   V5             8000 non-null     float64
13 |  7   V6             8000 non-null     float64
14 |  8   V7             8000 non-null     float64
15 |  9   V8             8000 non-null     float64
16 | 10  V9             8000 non-null     float64
17 | 11  V10            8000 non-null     float64
18 | 12  V11            8000 non-null     float64
19 | 13  V12            8000 non-null     float64
20 | 14  V13            8000 non-null     float64
21 | 15  V14            8000 non-null     float64
22 | 16  V15            8000 non-null     float64
23 | 17  V16            8000 non-null     float64
24 | 18  V17            8000 non-null     float64
25 | 19  V18            8000 non-null     float64
26 | 20  V19            8000 non-null     float64
27 | 21  V20            8000 non-null     float64
28 | 22  V21            8000 non-null     float64
29 | 23  V22            8000 non-null     float64
30 | 24  V23            8000 non-null     float64
31 | 25  V24            8000 non-null     float64
32 | 26  V25            8000 non-null     float64
33 | 27  V26            8000 non-null     float64
34 | 28  V27            8000 non-null     float64
35 | 29  V28            8000 non-null     float64
36 | 30  Amount         8000 non-null     float64
37 | 31  Class          8000 non-null     int64
38 | dtypes: float64(29), int64(3)
39 | memory usage: 2.0 MB


1 | # Count the occurrences of fraud and no fraud and print them
2 | occ = df['Class'].value_counts()
3 | print(occ)

1 | 0    7983
2 | 1      17
3 | Name: Class, dtype: int64

1 | # Print the ratio of fraud cases
2 | print(occ / len(df))

1 | 0    0.997875
2 | 1    0.002125
3 | Name: Class, dtype: float64
```

As you can see, the ratio of fraudulent transactions is very low. This is a case of class imbalance problem, and you're going to learn how to deal with this in the next exercises.



1.6 Plotting our data

From the previous exercise we know that the ratio of fraud to non-fraud observations is very low. You can do something about that, for example by **re-sampling** our data, which is explained in the next video.

In this exercise, you'll look at the data and **visualize the fraud to non-fraud ratio**. It is always a good starting point in your fraud analysis, to look at your data first, before you make any changes to it.

Moreover, when talking to your colleagues, a picture often makes it very clear that we're dealing with heavily imbalanced data. Let's create a plot to visualize the ratio fraud to non-fraud data points on the dataset `df`.

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3
4
5 def prep_data(df):
6     X = df.iloc[:, 1:29]
7     X = np.array(X).astype(float)
8     y = df.iloc[:, 29]
9     y = np.array(y).astype(float)
10    return X, y
11
12
13 # Define a function to create a scatter plot of our data and labels
14 def plot_data(X, y):
15     plt.scatter(X[y == 0, 0], X[y == 0, 1], label="Class #0", alpha=0.5, linewidth=0.15)
16     plt.scatter(X[y == 1, 0], X[y == 1, 1], label="Class #1", alpha=0.5, linewidth=0.15, c='
    r')
17     plt.legend()
18     return plt.show()

```

By visualizing your data you can immediately see how our fraud cases are scattered over our data, and how few are cases we have. A picture often makes the imbalance problem often very clear. In the next exercises we'll visually explore how to improve our fraud to non-fraud balance.

References